

Writeup for first project of
CMSC 420: “Data Structures”
Section 0101, Fall 2019

Theme: Heaps and Priority Queues

On-time deadline: Monday, 09-09, 11:59pm
(midnight)

Late deadline (30% penalty): Wednesday, 09-11,
11:59pm (midnight)

1 Overview

In this “warmup” programming project, you will remind yourselves of two basic, yet very powerful, data structures that you have talked about in courses such as CMSC132 and CMSC351: Binary **(Min-)Heaps** and **Priority Queues**. You will have to implement four classes:

- `LinkedMinHeap.java`
- `ArrayMinHeap.java`
- `LinearPriorityQueue.java`
- `MinHeapPriorityQueue.java`

You will first need to pull our starter code from [GitHub](#). This will check out **both** skeleton source code for your classes **and** documentation of your classes. You should first study the documentation of the classes under ‘doc’ to understand which operations you have to implement and how they should behave for various input configurations. The source code you will have to fill, alongside some starting tests, can be found under `projects.pqueue`.

Specifically, the first two classes under the (sub-)package `heaps` and the latter two under the (sub-)package `priorityqueues`. You should **not** move those classes from those packages! If you do, the submit server unit tests will **not** run against your implementation! We also provide the implementations and some unit testing suites for several simple data structures, as well as some custom `Exceptions` which are *checked* by various methods of `PriorityQueue`. You can consult the documentation and source code for more.

All the code that you write *can* be contained in those 4 source files. However, **this does not necessarily need to be the case!** For example, if you want, you can split your unit tests in multiple files or have some classes used by those four classes be defined in separate files. It's your code, your choice. The only thing that matters for the submit server is that the four `public` classes that you implement reside **exactly** where they are given to you in the code base's package structure.

2 Prerequisites

All prerequisites for handling this project are reasonable for an advanced CS student. We expect that you are well-familiar with **Binary Search Trees, Stacks, Lists, FIFO queues** and programming in Java. Skills harnessed by a typical UMD freshman course such as CMSC 131 / 132 are more than sufficient. You will need to remind yourselves of what a **binary heap** is and how insertions and min-deletions work, as well as the ways in which they can be represented in computer memory. The structure and operation of heaps are briefly touched upon in section 3.

In this class, we do not aim to test your programming, OOP or Java expertise. We want to teach you Data Structures. However, **a minimal amount of familiarity with certain OOP / Java constructs will be required.** You would probably benefit from solving the 3rd question from the 1st homework **before** you tackle this project, so that you are certain that you understand the basics of `Iterators` and checked `Exceptions`. You should also be familiar with what an interface is, what an anonymous inner class is, how it differs from a *functional* interface, etc

3 Short reminder on Heaps and Priority Queues

Heaps and Priority Queues are examples of data structures that are considered **background knowledge** for CMSC420. In this section, we will very briefly remind you of how heaps work and (re-)introduce Priority Queues.

3.1 Heaps

A **heap** is a [complete](#) binary tree (not a complete binary **search** tree!) where insertions always occur at the “rightmost” available position at the leaf level (the last position in a

level-order traversal of the heap). Furthermore, when considering deletions of elements, it is only possible to delete the element at the **root** of the tree. The key property of the heap is that every node's subtree contains elements whose values are **greater than or equal** to the node's element itself. Formally, this is known as a **minheap**; a **maxheap** is a heap where every node's subtree contains elements whose values are **smaller than or equal** to the node's element itself. It should then be clear that heaps are binary trees, but not binary search trees! :) In this project, we only care about **minheaps**.

After inserting a new element in the minheap, we have to make sure the heap property is maintained. To do this, we *percolate* the new element **upwards** as far as we can: for every parent, we need to make sure that the parent's value is smaller than or equal to our own. If not, we swap ("percolate upwards") the parent with the currently examined node. This process of "upwards percolating" *globally* guarantees the heap property, a fact that we can formally prove via induction.

For example, figure 1 shows the process of "upwards percolation" for a minheap after a new insertion. Figure 2 shows another example, where the process goes all the way up to the root!




Figure 1: Inserting 8 in a minheap. Since $10 > 8$, we need to swap 10 with 8. No further percolation is necessary, since $3 < 8$.




Figure 2: Inserting 2 in the heap of Figure 1 leads to consecutive swappings with 4 and 3, for a two-level upwards percolation.

As mentioned above, heaps only allow for the deletion of the root element. For minheaps, since the root element is also the **minimum element** (might not be unique, but that's fine), the operation is often referred to as "deleting the minimum" (**deleteMin**). In

the case of `deleteMin`, we retrieve the minimum element from the root node and replace it with the element in the **rightmost leaf** (the “last” element in a level-order traversal of the heap), which we subsequently throw away from the heap. Since `deleteMin` needs access to this “last element” of the heap, it is important, for efficiency reasons, to **always maintain a reference to this element**. After we copy over this element, we percolate it **down**, swapping the parent with the **minimum** of our two children at every level (figure 3). We can once again prove inductively that this “downwards percolation” process *globally* guarantees the heap property.




Figure 3: Deleting the minimum, replacing with “last” leaf element and percolating down.

Since heaps are complete binary trees, they can be implemented **very efficiently and compactly** using an **array** (figure 4). This is based on a breadth-first (level-order) enumeration of the nodes in the heap. This enumeration is exemplified in figure 4. Note that the node at index i has children at indices $2i + 1$ and $2i + 2$ (one or two children **might not even exist**, of course), whereas the parent of index i (**if existing**) is at index $\lfloor \frac{i-1}{2} \rfloor$.




Figure 4: Implementation of the resulting minheap of Figure 3 as an array.

Implementing your minheap in this way will be your task in `ArrayMinHeap.java`.

You will have to think about how your implementation will have to change from the linked implementation of `LinkedMinHeap.java` and how your percolation routines will need to be adjusted. It might help you to realize that **the exact same unit tests that you write for `LinkedMinHeap` should also be passable by `ArrayMinHeap`, and vice versa!** That is, if you were to just **copy and paste** all of the unit tests for one of the two classes in a separate file and test them against the other class, the only thing that you **should** have to change is the type of the reference in your `jUnit` tests from, say, a `LinkedMinHeap` to an `ArrayMinHeap` reference type.

3.2 Priority Queues

The Priority Queue is an Abstract Data Type (ADT) with a very simple property: Every element to be enqueued is attached a certain positive integer priority, which predetermines its **order** in the queue. By convention, **smaller integers are considered “higher” in terms of priority**, such that, for example, priority 1 is considered a higher priority than 3. Dequeueings only happen from the top of the queue, after which the element “before” the first one will be available for immediate processing. We see these kinds of queues all the time in real life. For example, people with certain disabilities and veterans are given higher priority when boarding a plane!

As an example, let’s consider two separate cases of enqueueings:




Figure 5: A simple FIFO queue enqueueing scenario. The person in the top can only be enqueued in the “back” of the queue.

In figure 5, we see an example of a traditional **FIFO (First In - First Out)** queue. In such a queue, we can only insert at the end, so **longer queues have longer average wait times, irrespective of the properties of the elements that are processed**. This rigidity is unacceptable in practice, and Priority Queues are used almost universally. A perfect example is **Operating Systems**, which maintain a Priority Queue of processes that are to be allocated CPU time. Some processes are so much more important than others that they **must** be given higher priority!

Despite their apparent simplicity, FIFO queues can be implemented in a number of different ways, each with its own different properties. For example, some implementations allow for faster enqueueings than dequeuings (or vice versa), while others use contiguous or dis-contiguous memory! Please feel free to examine the package `fifoqueue` for three different implementations of FIFO queues, **all of which adhere to the common interface**, but use **different implementations** with **wildly different properties, efficiency characteristics and memory footprints!**

In contrast with a traditional FIFO Queue, when considering a Priority Queue, every element has a certain **positive** priority attached to it:




Figure 6: A simple Priority Queue enqueueing scenario.

As you are even reading these lines and looking at Figure 6 you are probably envisioning a **linear** implementation of a Priority Queue. **This is good**; keep these visualizations in mind when implementing `LinearPriorityQueue`. Naturally, such an implementation of a Priority Queue would be **very inefficient**, since it leads to $\mathcal{O}(n)$ enqueueings and, depending on the underlying implementation, maybe even $\mathcal{O}(n)$ dequeuings! A much better way to implement Priority Queues is a **binary minheap**, and this will be your task when implementing the class `MinHeapPriorityQueue`. Since minheaps are **complete** binary trees, and complete binary trees are also **balanced** (except for maybe one level of imbalance), we have enqueueings and dequeuings in $\mathcal{O}(\log_2 n)$. These efficiency **guarantees** of the minheap make it a **very desirable choice for implementing Priority Queues**.

4 Tips / Hints / Guidelines

The following are just some guidelines that **may** be helpful for you to implement your project.

- Familiarize yourselves with the most elementary `git` commands (`clone`, `pull`, `add`, `commit`, `push`, `alias`). Our ELMS page has startup resources for you under a page called “More Resources”. Your elementary `git` knowledge will also be tested on your **first homework**. We **highly** suggest that you **pull** the code (this is literally a single

command) **instead of downloading the zip** from GitHub. This is because, if we ever post a correction on the docs or the source code, you will only need to type **one single command** to pull **only the changes themselves**, without affecting your existing solutions! If you always download the ZIP whenever a correction is made, you will have to nitpick the files that you copy-and-paste and **possibly overwrite your work by mistake**. In fact, this is **guaranteed to happen** if you download the ZIP the first time, work on the project for a while, and when a correction is posted you decide to try out `git` and do a `git clone`. This will **immediately overwrite everything in your current directory and replace it with the file state of the latest upstream commit**. **IN THEIR GRADUATE STUDENT CAREER, THE INSTRUCTOR LOST TWO WEEKS' WORTH OF CODE BECAUSE OF THIS MISTAKE.**

- **Read the JavaDocs** of the interfaces that you will have to extend (`MinHeap.java`, `PriorityQueue.java`). Opening `index.html` from your favorite web browser and then navigating around is a fine way to start.
- **TEST FIRST, IMPLEMENT SECOND. WE CANNOT STRESS HOW EASIER THIS WILL MAKE YOUR DEVELOPMENT.**
- Read the provided data structures' source code **and** unit testing frameworks for ideas on implementation and testing.
- Think about how you would write a **stress** test for a method. You want to make sure that the method passes **any conceivable** input (of the appropriate type, of course), where **“not passing”** means **“throwing any conceivable error”**.
- Read the source of the existing `StudentTests` to see how you can write a unit test that **expects** a certain error to be thrown!
- Consider switching from Eclipse to an IDE such as [NetBeans](#) or [IntelliJ](#). We provide IntelliJ resources under “Resources” on ELMS.
- For this course, you will **not** need to play around with functional programming features and **Streams**, first introduced in Java 8. However, `IntStreams`, `Ranges`, lambdas and `Collectors` can make your unit tests **much easier** to write. For example, here are two lines of code from one of our unit tests:

```
List<Integer> priorities = IntStream.range(1, MAX_PRIORITY + 1).boxed().  
    collect(Collectors.toList());  
Collections.shuffle(priorities, r);
```

While one could quite easily write the two lines above using for-loops, the entire point behind most functional Java 8 features is to make developers [avoid tedious, non-parallelizable and error-prone loops](#).

5 Submission / Grading

Credit in this project is defined by the number of [submit server](#) unit tests that you pass. However, **we will also be inspecting your source code to make sure that your work adheres to certain standards.** For example, the `ArrayMinHeap` class needs to be implemented using contiguous storage (e.g `ArrayList`, `Vector`, raw arrays), while your implementation of `MinHeapPriorityQueue` **must** use either one of your `MinHeap` implementations! **In order to get any credit for the unit tests that correspond to a specific class, your implementation should adhere to the documentation's guidelines.** In particular, using built-in Java primitives or third-party libraries that implement `MinHeaps` and `PriorityQueues` are a **very easy way to ensure zero credit!** Those requirements are also repeated in the `JavaDocs`.

Projects in this class are different from your typical 131/2 projects in that **we do not maintain an Eclipse - accessed CVS repository** for you or us. This means that you can no longer use the Eclipse Course Management Plugin to submit your project on the submit server. This turns out to be a good thing, since it frees you up from the need to use Eclipse if you don't want to.

To submit your project, run the script `src/Archiver.java` as a **Java application** from your IDE (tested with Eclipse and IntelliJ). This will create a **.zip archive** of your **entire project directory** at the same directory level of the project folder (so, “next” to it), **without** including the hidden directory `.git` and `doc`, which are useless to the server for testing purposes. For example, if your project directory is under `/home/users/me/mycode/project1/`, this script will create the .zip archive `/home/users/me/mycode/project1.zip`, which will contain `src`, and any other directories that you may have, but will **not** contain the directories `.git` and `doc`. After you have done this, upload the archive on the submit server as seen on figure 7.

[illegible]

Making another submission

- use automatic submission tools,
- edit and submit code in the browser (discouraged)
- upload source files

Uploading a submission

You can submit [a zip file](#) or [multiple text files](#).

file(s) for submission

File(s) to Submit: Choose File No file chosen

Figure 7: Uploading your project on the submit server.

All tests are release tests, and you can submit **up to 5 times** every 24 hours. We urge you to **unit-test your code thoroughly before submitting**: treat every token like a bar of gold that is not to be wasted! We will **not** share the source code of the unit

tests with you, not even after the deadline for the project!

We maintain your **highest-scoring submission** for grading purposes. Finally, for the late deadline, we take 30% off your maximum possible score. This means that, if you submit late, passing all the unit tests will give you 70% of the total grade.

Finally, we should remind you that for the past few years the [Software Similarity Detection System MoSS](#) has been incorporated into the CS department's submit server. For n student submissions, it is **ridiculously easy** (literally a single click) for us to run MoSS against all $\binom{n}{2}$ pairs of submissions. MoSS is tuned towards higher than 50% Recall, which means that plagiarized submissions **will** be caught. We would much rather be spending time teaching you data structures and assisting you with your queries than going back and forth with the Honor Council; **help us help you!**